

Neural Networks and Deep Learning

CS5242

Wei WANG

National University of Singapore

cs5242@comp.nus.edu.sg

Recap



- AI, ML, and DL
- Topics: MLP, CNN, RNN
- Pre-requisite and workload
 - Run assignment 0
 - Understand house price prediction

Recap

- House price prediction
 - **Sample** (example): (one house, price) is a sample
 - **Feature** $x \in R^m$: attributes used to represent the house, e.g. size, #floors
 - Ground-truth **label** y : the real price
 - Model: linear regression, $\tilde{y} = w^T x + b, w \in R^m, b \in R$
 - **Parameters**: w, b
 - **Prediction** \tilde{y} : the price predicted by our ML model
 - **Loss function**: objective for training the model.
 - $J(w, b) = \frac{\sum_{\langle x, y \rangle \in S_{train}} L(x, y | w, b)}{|S_{train}|} = \frac{\sum_{i=1}^n L(x^{(i)}, y^{(i)} | w, b)}{n}$
 - $L(x, y | w, b) = |\tilde{y} - y|^2$
 - Back-propagation and gradient descent

Training by gradient descent

- Gradient descent (GD) algorithm for optimization



α is called the learning rate, which controls the moving step length. It is important for convergence. If it is large, w would oscillate around the optimal position. If it is small, it would take many iterations to reach the optimal position.

Initialize w as w_0

Compute $\frac{\partial J}{\partial w_0}$, negative;

Move w from w_0 to the right by

$$w_1 = w_0 - \alpha \frac{\partial J}{\partial w_0}$$

Compute $\frac{\partial J}{\partial w_1}$, negative;

Move w from w_1 to the right by

$$w_2 = w_1 - \alpha \frac{\partial J}{\partial w_1}$$

Compute $\frac{\partial J}{\partial w_2}$, positive


Move w from w_2 to the left by

$$w_3 = w_2 - \alpha \frac{\partial J}{\partial w_2}$$


Gradually decrease J and move w to the optimal position

Back-propagation

- Forward
- $X \in R^{m*n}, Y \in R^{1*n}$
- $V = w^T X, \in R^{1*n}$
- $\tilde{Y} = V + b, \in R^{1*n}$
- $E = \tilde{Y} - Y, \in R^{1*n}$
- $L = E^2, \in R^{1*n}$
- $J = \text{numpy.average}(L) \in R^+$



Back-propagation



• Backward

- $\frac{\partial J}{\partial L} = \left[\frac{1}{n}, \frac{1}{n}, \frac{1}{n}, \dots \right], \in R^{1 \times n}$
- $\frac{\partial J}{\partial E} = \frac{\partial J}{\partial L} \times \frac{\partial L}{\partial E} = \frac{\partial J}{\partial L} \times 2E = 2E/n, \in R^{1 \times n}$
- $\frac{\partial J}{\partial \tilde{Y}} = \frac{\partial J}{\partial E} \times \frac{\partial E}{\partial \tilde{Y}} = \frac{\partial L}{\partial E} \times [1, 1, 1, \dots] = 2E/n, \in R^{1 \times n}$
- $\frac{\partial J}{\partial b} = \frac{\partial J}{\partial \tilde{Y}} \cdot \frac{\partial \tilde{Y}}{\partial b} = \frac{\partial L}{\partial \tilde{Y}} \cdot [1, 1, 1, \dots], \in R$ (dot product)
- $\frac{\partial J}{\partial V} = \frac{\partial J}{\partial \tilde{Y}} \times \frac{\partial \tilde{Y}}{\partial V} = \frac{\partial L}{\partial \tilde{Y}} \times [1, 1, 1, \dots] = 2E/n, \in R^{1 \times n}$
- $\frac{\partial J}{\partial w} = \left(\frac{\partial L}{\partial V} \frac{\partial V}{\partial w} \right)^T = X \left(\frac{\partial L}{\partial V} \right)^T, \in R^m$ (matrix-matrix product)

\times : element-wise multiplication

Element-wise multiplication?
dot product?
matrix product?
transpose ?






Shape check: for every node in the graph, its shape should be the same during forward and backward.



Gradient descent

<http://ruder.io/optimizing-gradient-descent/>

Loss contour




Source from
<https://goo.gl/ULkt2Y>






Data normalization

- The attributes of a house varies in scale
- Normalize the attributes into similar scale



Challenges of gradient descent

- Local optimum
- Saddle points



Gradient descent (GD)

Stochastic gradient descent (SGD)

Mini-batch SGD

Mini-batch SGD with momentum

RMSProp

Adam

The Evolution of Gradient Descent

- <https://www.youtube.com/watch?v=nhqo0u1a6fw>

Bias and variance

Training and testing

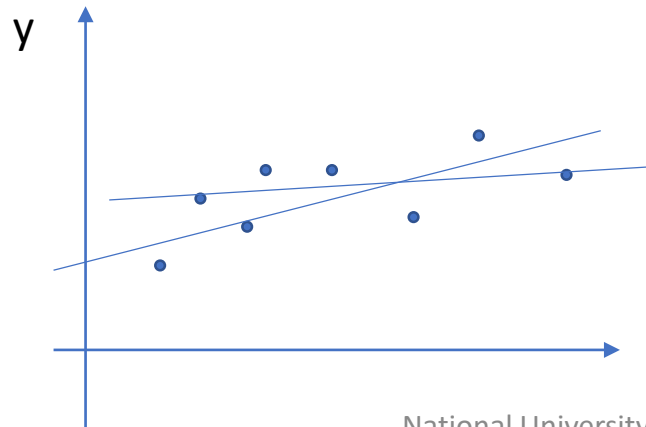
- Train a model over experience data (training samples); and then
- Deploy the model to do prediction for **new** data;

Training and testing performance

- Training
 - minimize the error over training data
- Testing
 - Fix the model parameters
 - Make predictions on unseen data samples (i.e. new data)
- Ideal case
 - Small training error and small testing error.


Underfitting

- For example,
 - $\tilde{y} = xw + b$
- The model is too simple to model the data
 - Low capacity/complexity

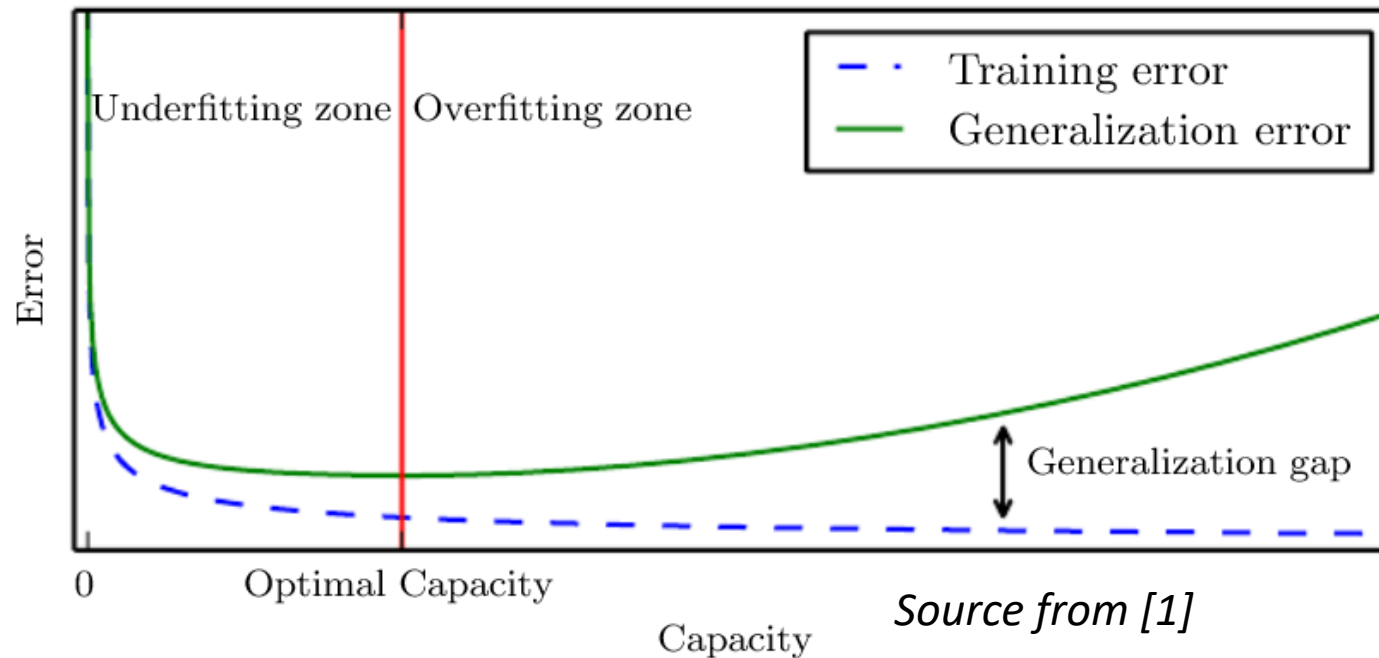


Overfitting

- Good performance on seen data, i.e. training data
- Bad performance on unseen data, i.e. test data
- $\hat{f}(x) = f(x) + \varepsilon$
- We learn the model $\hat{f}(x)$ over noisy data $f(x) + \varepsilon = y$
 - Small training error \rightarrow the model fits the noise very well
 - It may fail to fit the new data (not noise) \rightarrow large test error



Underfitting and overfitting



Bias and variance

$$\mathbb{E}[(y - \hat{f}(x))^2] = \text{Bias}[\hat{f}(x)]^2 + \text{Var}[\hat{f}(x)] + \sigma^2$$

where

$$\text{Bias}[\hat{f}(x)] = \mathbb{E}[\hat{f}(x) - f(x)]$$


and

$$\text{Var}[\hat{f}(x)] = \mathbb{E}[\hat{f}(x)^2] - \mathbb{E}[\hat{f}(x)]^2$$


- <https://ml.berkeley.edu/blog/2017/07/13/tutorial-4/>

Bias and variance

- <https://elitedatascience.com/bias-variance-tradeoff>



All in one picture



Source from [1]

Diagnosis of underfitting and overfitting


- $\tilde{y} = w_1x + b$
- $\tilde{y} = w_1x + w_2x^2 + b$
- $\tilde{y} = w_1x + w_2x^2 + w_3x^3 + b$
- $\tilde{y} = w_1x + w_2x^2 + w_3x^3 + w_4x^4 + b$
- Solution
 - Compare training error and testing error
 - Select the optimal capacity
- This process is like training a model over the test data

Diagnosis of underfitting and overfitting

- Validation data
- Benchmark overfitting
 - ImageNet

Solutions for underfitting

- Increase model capacity
 - Capacity?
 - Use more features
 - Use complex models
 - $\tilde{y} = w_1x + w_2x^2 + w_3x^3 + w_4x^4 + w_5x^5 + b$



Solutions for overfitting

- Decrease the gap |training error – test error|
 - Training data
 - See almost all cases although not every sample in the test data
 - Capacity
 - Regularization of the hypothesis space
 - Add prior or constraint to the functions.
 - L2
 - Dropout
 - Parameter sharing
 - Ensemble
 - Dropout
 - Early stopping

L2 regularization

- $\min_{\theta} J(\theta)$

- $\min_{w,b} J(w, b) = \frac{\sum_{\langle x, y \rangle \in S_{train}} L(x, y | w, b)}{|S_{train}|} = \frac{\sum_{i=1}^n L(x^{(i)}, y^{(i)} | w, b)}{n}$

- $\min_{\theta} J(\theta) + \lambda |\theta|^2$

- Partial derivative

- derivate from both $J(\theta)$ and the regularization term

- $\frac{\partial J}{\partial \theta} + 2\lambda \theta$

- Plot $J(\theta) + \lambda |\theta|^2$ instead of $J(\theta)$

- Real experience

- https://github.com/BVLC/caffe/blob/master/examples/cifar10/cifar10_full.prototxt#L140

- Bias and variance
- <https://www.youtube.com/watch?v=SjQyLhQIXSM>

Reference

- [1] Goodfellow Ian, Bengio Yoshua, Courville Aaron. Deep learning. MIT Press. <http://www.deeplearningbook.org>. Chapter 5.
- <https://www.analyticsvidhya.com/blog/2017/03/introduction-to-gradient-descent-algorithm-along-its-variants/>
- <https://elitedatascience.com/bias-variance-tradeoff>
- Curse of dimensionality <https://goo.gl/4UT253>